

# Sort it Like You Mean It: Discovering Semantically Interesting Attribute Augmentations to Sort Tables

Akash Khatri

University of Utah

akash.m.khatri@utah.edu

Mir Mahathir Mohammad

University of Utah

mahathir.mohammad@utah.edu

El Kindi Rezig

University of Utah

elkindi.rezig@utah.edu

## ABSTRACT

Sorting is a fundamental operation in table analysis. Data scientists frequently sort tables to uncover key insights—for example, identifying the top 10 products by sales. However, this process is largely manual. Data scientists must (1) understand the semantics behind the sorting they wish to apply, and (2) ensure the necessary attributes are present—often requiring manual augmentation of the table. But what if data scientists could receive suggestions for semantically meaningful ways to sort a table, powered by automatic augmentations from a data lake? In this demo, we present INSIGHTSORT, an end-to-end system that recommends attribute augmentations to enable richer, more insightful sorting for table exploration. INSIGHTSORT works by: (1) discovering potential augmentations by linking the input table with relevant data lake tables, and (2) leveraging a Large Language Model (LLM) to synthesize the top-k diverse sorting attributes based on their semantics. A companion video is available at [1].

## 1 INTRODUCTION

Sorting is a critical operation in both data management and exploratory analysis. For instance, when analyzing a table of employee records, sorting by attributes such as income or seniority can surface meaningful patterns or generate insightful reports. This might reveal, for example, employees who have stayed with the company for many years but have received relatively few raises—highlighting potential disparities or areas for further investigation.

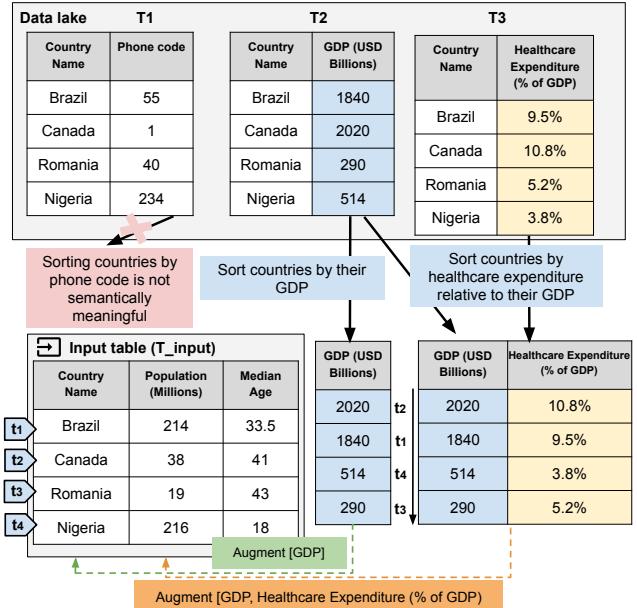
However, this process assumes that users already know which attributes to sort by and that these attributes are readily available in the table. It also requires users to understand the semantics of the desired sorting logic—for example, recognizing that sorting employees by years of seniority may involve computing tenure from hire dates.

Data lakes—whether organizational or public [8]—are becoming increasingly ubiquitous, offering opportunities for data augmentation for several tasks [5]. We propose an approach to identify candidate sorting attribute sets within a data lake to augment a given input table. This enables the table to be sorted in semantically meaningful ways that go beyond its original schema.

Even when users have a general sense of the sorting semantics they want to apply, distinguishing between subtle variations in sorting strategies can be challenging. For example, if users wish to sort countries by “economic performance,” potential sorting attributes might include GDP, GDP per capita, GDP growth rate, or inflation rate. Each of these captures a different facet of economic performance, and identifying the most relevant attributes often requires significant manual effort.

What if users could receive recommendations for interesting ways to sort a table—even when the required attributes are not

initially present? These attributes could be automatically sourced from a data lake or even generated using a Large Language Model (LLM) [2].



**Figure 1:** The input table can be augmented with several attributes from the data lake to enable semantically interesting sorting of countries. The tuple IDs are shown next to the suggested augmentation attributes to reflect a descending sort.

*Example 1.1.* A data scientist, Lou, has an input table of countries with the Population and Median Age attributes (Figure 1). Lou also has access to a data lake containing numerous tables, but it’s unclear which sets of attributes can augment the current table to enable insightful sorting. The data lake tables suitable for augmentation must be joinable with the input table. For instance, joining  $T_1$  with the input table using the columns Country Name introduces a new attribute GDP, enabling countries to be sorted by GDP, a crucial measure of economic performance.

Another semantically meaningful way to sort the input table is by joining it first with  $T_2$  and then  $T_3$ , adding two attributes: GDP and Healthcare Expenditure. These attributes allow for sorting the countries first by GDP and then by healthcare expenditure, highlighting how much countries invest in healthcare relative to

their GDP. However, not all sorting attributes are useful. For instance, Sorting countries using their phone code is not semantically meaningful.

We present a demo of **INSIGHTSORT**, an end-to-end system that (1) creates indexes that keep track of the sorting columns in data lake tables; (2) efficiently determines, given an input table, what data lake columns are joinable with it; (3) selects sorting attributes that are semantically meaningful; and (4) diversifies the sorting strategies based on their semantic interestingness. Finally, evaluating LLMs for table sorting is an unexplored area. This demo offers VLDB participants a unique opportunity to assess the performance of several leading LLMs—such as GPT-4o [2], Perplexity AI [9], and Claude [3]—on this important table analysis task.

**Related work.** Several techniques for join discovery have been proposed [4, 11]. In this line of work, users specify a query (e.g., example table) and the goal is to find joinable tables from the data lake. In Metam [5], the goal is to perform attribute augmentation for a given downstream task by specifying its utility function. **INSIGHT-SORT** differs from existing work because (1) its goal is orthogonal to the join discovery being used, one can choose any join discovery method to find joinable tables; and (2) it searches for semantically-meaningful sorting attributes to augment tables, which is a target task that has never been explored in prior work.

## 2 SYSTEM OVERVIEW

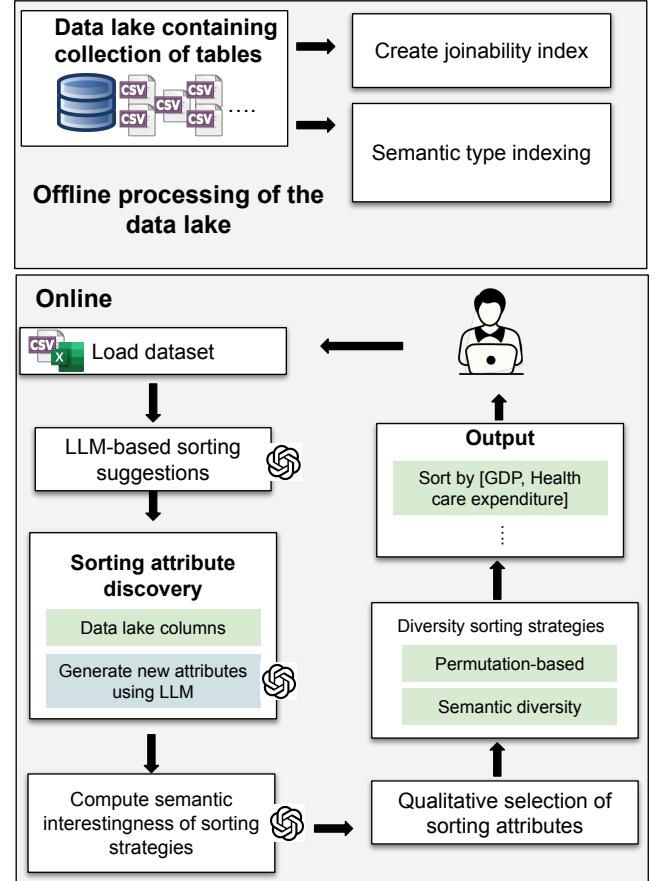
Figure 2 illustrates the high-level workflow of **INSIGHTSORT**. In a nutshell, **INSIGHTSORT** ingests a data lake of tables and builds necessary indexes to navigate it in the offline phase. In the online phase, **INSIGHTSORT** takes as input a table. **INSIGHTSORT** finds possible attribute augmentations for the input table from the data lake, scores them by their semantic interestingness, and finally diversifies the sorting results.

### 2.1 Offline processing of the data lake

In this phase, **INSIGHTSORT** ingests a data lake of tables and indexes it for joinability testing, i.e., given an input column, which data lake columns are joinable with it.

**2.1.1 Building the joinability index.** Since the data lake can contain a large number of tables, it is important to build indexes that enable efficient joinability testing. To this end, **INSIGHTSORT** proceeds as follows:

- (1) We generate vector embeddings for each column in the data lake by following the approach proposed in [4], which leverages the pre-trained language model SBERT [10]. Briefly, this method constructs a textual summary for each column—such as a list of unique values—and then uses the language model to encode these summaries into vector embeddings.
- (2) After the vector embeddings of all columns have been created, **INSIGHTSORT** index them using Hierarchical Navigable Small World (HNSW) [7]. This index enables fast Approximate Nearest Neighbor (ANN) queries to check if a query column embedding matches a data lake column embedding within a distance threshold.



**Figure 2: High-level workflow of **INSIGHTSORT**.** **INSIGHTSORT** ingests and indexes a data lake of tables (offline processing), which is then used in the online phase to discover sorting attributes.

**2.1.2 Semantic types indexing.** **INSIGHTSORT** creates another HNSW index that indexes the vector embeddings of textual description of columns. To this end, **INSIGHTSORT** performs the following:

- (1) Because column names are most often not descriptive, for each table, **INSIGHTSORT** prompts the LLM (GPT-4o [2]) with the attribute names and sample values, and requests the semantic types of each column. The semantic type of a column is its textual description. For instance, a column’s semantic type can be “US universities” if it contains names of US universities.
- (2) **INSIGHTSORT** uses pre-trained language models (sentence transformers [10]) to generate vector embeddings of the semantic types.
- (3) Finally, **INSIGHTSORT** indexes those vector embeddings using an HNSW index for fast ANN query answering.

### 2.2 Online phase

In the online phase, **INSIGHTSORT** leverages the previously created data lake indexes to find joinable columns to a given query table, and possible join columns for a given sorting description.

**2.2.1 Uploading the input table.** INSIGHTSORT allows users to upload their input table. INSIGHTSORT will then look for sorting columns by augmenting it with data lake columns or creating new columns using the LLM.

**2.2.2 LLM-based sorting suggestions.** INSIGHTSORT prompts the LLM with the input table schema, and sample rows, and asks for “interesting” sorting attributes that can be added to the input table. The LLM identifies a list of attributes that can be added to sort the rows. For instance, in the example table  $T_2$  in Figure 1, GPT-4o [2] suggested the sorting attributes in Table 1 (we are only including a sample of them).

**2.2.3 Sorting attribute discovery.** Given the suggested sorting attributes by the LLM, INSIGHTSORT attempts to locate them in the data lake. Of course, those attributes will not necessarily have the same name as suggested by the LLM, so INSIGHTSORT proceeds as follows:

- (1) **Vector embedding generation:** INSIGHTSORT generates vector embeddings of the suggested sorting attribute textual description using a pre-trained language model. INSIGHTSORT uses Sentence Bert [10] to produce those embeddings.
- (2) **Locating the query attributes:** INSIGHTSORT finds the data lake columns that match the query column by performing an ANN query on the semantic type HNSW index. After this, INSIGHTSORT has to determine if the tables of those matching columns are joinable with the query table. To this end, it uses the joinability index that was created offline.

Additionally, INSIGHTSORT can also ask the LLM to generate a sorting attribute directly. This is appropriate for cases where the sorting attribute is “general-knowledge” and not private (e.g., company-specific data).

Attribute	Description / Use for Sorting
Population (millions)	Sort by population size
GDP per Capita (USD)	Economic comparison across countries

**Table 1: A sample of the suggested sorting attributes by GPT-4o for table  $T_2$  in Figure 1.**

**2.2.4 Computing semantic interestingness of sorting columns.** We now define the semantic interestingness of a set of sorting attributes with respect to a query table:

**Semantic interestingness of a sorting attribute.** Given an input table  $T$  and a candidate set of columns  $C = \{c_1, c_2, \dots, c_m\}$ , the semantic interestingness of a given sort strategy (attributes by which to sort) captures the following two components:

- **Semantic utility.** How semantically interesting the attributes in  $C$  are in generating an interesting sort order for  $T$  ( $\text{semantic\_utility}(T, C)$ ). This is computed by prompting an LLM with the column names of  $T$  and  $C$  and asking it to derive a score for how interesting sorting  $T$  with  $C$  is.
- **Column uniqueness.** We also want to make sure the attributes in  $C$  do not overlap significantly with the attributes in  $T$ . That is, we don’t want to augment the table with duplicate attributes.

Column uniqueness is computed by averaging the cosine similarity between the columns in  $C$  and the attributes in  $T$  (we get the top-1 match from  $T$  for each column in  $C$ ).

In summary, the semantic interestingness is defined as:

$$I(T, C) = \frac{\text{semantic\_utility}(T, C) + \text{uniqueness}(T, C)}{2} \quad (1)$$

Seamntic interestingness allows us to determine, for each joinable table in the data lake, whether adding its attributes to the input table would result in semantically-interesting sorting attributes.

**2.2.5 Diversifying the results.** INSIGHTSORT returns a set of *diverse* sorting columns. That is, we want to make sure we minimize duplicate sorting attributes. To this end, we adopt two strategies:

- **Permutation-based:** for two given sorting attributes  $A$  and  $B$ , applied on the input table  $T$  with rows identifiers  $\{1, 2, \dots, n\}$ , result in row permutations  $S_A$  and  $S_B$ . That is,  $A$  and  $B$  result in different shuffling of the input table’s rows. Because we want to produce diverse sorted tables, this means that two sorted tables should have significantly different permutations of rows. To this end, we use the hamming distance or levenshtein distance between the permutation arrays of two sorted tables to determine if they are diverse if the distance is greater than  $\theta$ . Presumably, if two sorting attributes are correlated (e.g., GDP and HDI), they are likely to result in the same permutation of rows.
- **Semantic diversity:** In this strategy, we want to diversify sorting attributes if they are “semantically” diverse. To this end, INSIGHTSORT (1) generates the textual description of each sorting strategy using the LLM; (2) generates vector embeddings of those descriptions using a pre-trained language model (e.g., sentence transformers); and (3) uses an HNSW index to determine if two sorting strategies are semantically similar, i.e., if they are within a given distance threshold  $d$ .

### 3 DEMONSTRATION PLAN

We demonstrate INSIGHTSORT on the following real-world datasets: (1) Stackoverflow survey dataset <sup>1</sup>; (2) Adult dataset <sup>2</sup>; and (3) Chicago crime dataset <sup>3</sup>. As for the data lakes, we will be using a repository of over 1000 CSV files that we crawled from data.gov.

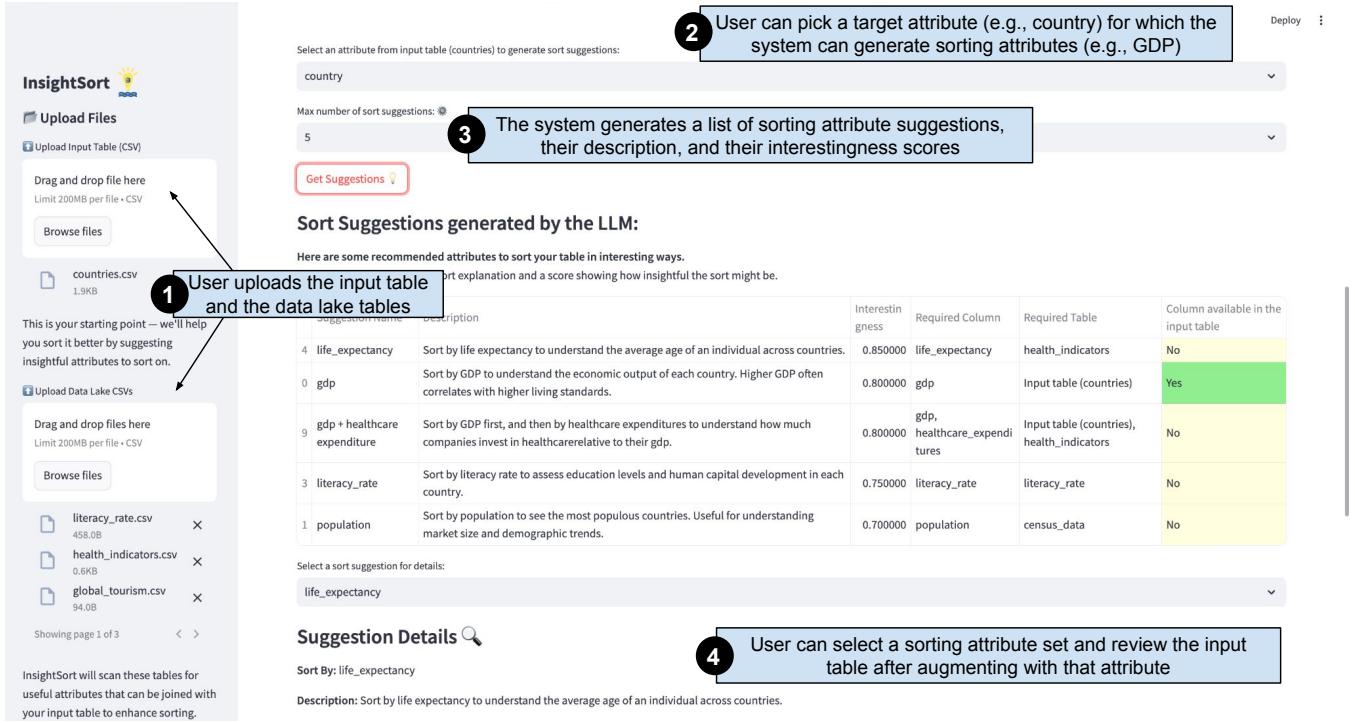
**Demonstration outline.** In this demo, we aim to showcase (1) the interactive interface of INSIGHTSORT that allows users to explore various table augmentations for different sorting semantics; (2) the importance of interestingness scores in exposing various “interesting” sorting attribute sets; and (3) the system’s diversity sorting which allows users to inspect sorting attributes that are semantically unique. We now present the demonstration scenario:

- ① **Uploading the input table and the data lake.** The participants start by uploading their input table as well as the table repository from which INSIGHTSORT will discover augmentations (Figure 3 ①).
- ② **Customizing joinability thresholds.** Once the data lake has been uploaded, INSIGHTSORT finds joinable columns with the input

<sup>1</sup><https://survey.stackoverflow.co>

<sup>2</sup><https://www.kaggle.com/datasets/wenrui/adult-income-dataset>

<sup>3</sup><https://www.chicago.gov/city/en/dataset/crime.html>



**Figure 3: The user interface of INSIGHTSORT. Participants will have the opportunity to discover several sorting attribute sets with different semantics.**

table. Participants can adjust the joinability thresholds that are internally used by INSIGHTSORT to determine if two columns are joinable. The goal here is to show that less joinable columns can create several spurious augmentations.

③ **Picking a target attribute.** INSIGHTSORT allows users to pick a particular target attribute (e.g., country) for which INSIGHTSORT will discover sorting attributes (e.g., GDP) as illustrated in Figure 3 ②. Users can also elect not to choose a target attribute, in which case INSIGHTSORT will discover sorting attributes for all possible attributes in the input table.

④ **Sorting attribute suggestions.** INSIGHTSORT discovers a list of sorting attributes from the data lake, and displays their description, as well as their interestingness scores to the user (Figure 3 ③).

⑤ **Input table augmentation.** Upon selecting a particular sorting attribute set suggestion, INSIGHTSORT proceeds to augment the input table with the selected attribute set (Figure 3 ④).

⑥ **Comparing sorting attributes.** INSIGHTSORT allows users to compare various sorting sets based on how unique they are. This uniqueness is determined by how similar the sorting permutes the tuples of the input table, and how semantically unique it is. Users can adjust the thresholds of diversity to explore different sorting suggestions. The goal is to demonstrate that including no diversity filtering results in a large amount of similar sorting suggestions.

⑦ **Comparing different LLMs.** INSIGHTSORT will also feature connections to LLMs other than GPT-4o. Those include Gemini [6], Perplexity AI [9], and [3]. This will allow the participants to select sorting suggestions across different LLMs, which would provide a vast array of options.

*Demonstration engagement.* In addition to a guided demonstration of INSIGHTSORT, we will encourage participants to upload their own datasets to perform an interactive table exploration session on them. Participants will be able to interactively examine sorting attributes that are semantically relevant to their task and compare them across different LLMs, joinability thresholds, and interestingness scores.

## REFERENCES

- [1] [n.d.]. InsightSort demo video. <https://drive.google.com/drive/folders/1Hke2MTKtGF0FRHmbOwx-nNIFPUcojC9?usp=sharing>
- [2] Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. 2023. Gpt-4 technical report. *arXiv:2303.08774* (2023).
- [3] Anthropic. 2023. Claude. <https://www.anthropic.com/index/introducing-claude> Constitutional AI-based large language model.
- [4] Yuyang Dong, Chuan Xiao, Takuma Nozawa, Masafumi Enomoto, and Masafumi Oyamada. 2023. DeepJoin: Joinable Table Discovery with Pre-Trained Language Models. *Proc. VLDB Endow.* 16, 10 (June 2023), 2458–2470. <https://doi.org/10.14778/3603581.3603587>
- [5] Sainyam Galhotra, Yue Gong, and Raul Castro Fernandez. 2023. METAM: Goal-Oriented Data Discovery. *arXiv preprint arXiv:2304.09068* unspecified, unspecified (2023), unspecified–unspecified.
- [6] Google DeepMind. 2023. Gemini. <https://deepmind.google/technologies/gemini/> Multimodal large language model.
- [7] Yu A. Malkov and D. A. Yashunin. 2020. Efficient and Robust Approximate Nearest Neighbor Search Using Hierarchical Navigable Small World Graphs. *IEEE Trans. Pattern Anal. Mach. Intell.* 42, 4 (April 2020), 824–836.
- [8] Fatemeh Nargesian, Erkang Zhu, Renée J. Miller, Ken Q. Pu, and Patricia C. Arocena. 2019. Data lake management: challenges and opportunities. *Proc. VLDB Endow.* 12, 12 (Aug. 2019), 1986–1989. <https://doi.org/10.14778/3352063.3352116>
- [9] Perplexity AI. 2023. Perplexity AI. <https://www.perplexity.ai>
- [10] N Reimers. 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. *arXiv preprint arXiv:1908.10084* (2019).
- [11] El Kindi Rezig, Anshul Bhandari, Anna Fariha, Benjamin Price, Allan Vanterpool, Vijay Gadepally, and Michael Stonebraker. 2021. DICE: data discovery by example. (2021).